

CO395 Machine Learning
CBC #2
Artificial Neural Networks

Group 1

Yong Wen Chua, ywc110
Thomas Morrison, tm1810
Marcin Baginski, mgb10
Marcin Kadziela, mk4910

Contents

Implementation Details	3
Overview	3
Description of the MATLAB functions	3
Evaluation	4
Clean dataset	4
Single six-output network	4
Six single-output networks	5
Noisy dataset	5
Single six-output network	5
Six single-output networks	6
Discussion of results	7
Questions	8
Optimal topology	8
Overfitting	11
Six-output vs. single-output networks	11
Ideal parameter optimisation	12
Code Flowchart	13

Implementation Details

Overview

This report describes the work conducted in the second assignment of the Machine Learning course. The main goal of this coursework was to investigate and learn how to use MATLAB's Neural Network Toolbox. The task that we were given was a typical pattern recognition and classification problem. We were supposed to train, test and evaluate one network with multiple outputs and multiple single-output neural networks (one for each class). We started this assignment by optimising parameters for each of these networks. The approach taken was optimising topology together with the training function first. In later phases we focused on other parameters (such as minimum gradient or goal) which could reduce the problem of overfitting. Eventually, we used the 10-fold cross-validation in order to compare the performance of the two network types.

Description of the MATLAB functions

- `nFoldValidate(examples, classifications, n, networkType)` - performs the n-fold cross-validation. Takes as arguments the examples, classifications, number of folds and the network type ('single' or 'multi'). Based on the network type, it either generates multiple single-output networks or one multi-output network and validates them. It returns a cell array which contains the confusion matrices for each fold
- `generateMultiOutputNetwork(xANN, yANN, layers, neurons, transferFcn, trainingFcn, lr, grad, goal)` - generates a multiple-output network based on the arguments which are passed to the function. These include the number of layers, neurons in a layer, training function, transfer function, learning rate, stopping gradient and the performance goal. It is used within `nFoldValidate()`
- `generateSingleOutputNetworks(xANN, yANN, layers, neurons, transferFcn, trainingFcn, lr, grad, goal)` - generates multiple single-output networks based on the same arguments as the previous function
- `testANN(network, inputs)` - produces a vector of predictions for the inputs using provided neural network (either the multi or single output one)
- `recallPrecisionF1(confusionMatrix)` - calculates the recall, precision and F_1 measures after taking the confusion matrix as an argument
- `plotAverageF1Measure(confusionMatrixSingle, confusionMatrixMulti)` - takes two arrays of confusion matrices (one for each fold) from the single and multi-output networks and plots the average F1 measure for each fold. Its output is a direct answer to the question from Part VIII
- `confusion_matrix(actual, predicted, possible_outcomes)` - calculates the confusion matrix given a column vector of actual and predicted outcomes

Evaluation

The average F_1 measure per fold for both types of networks is presented below:

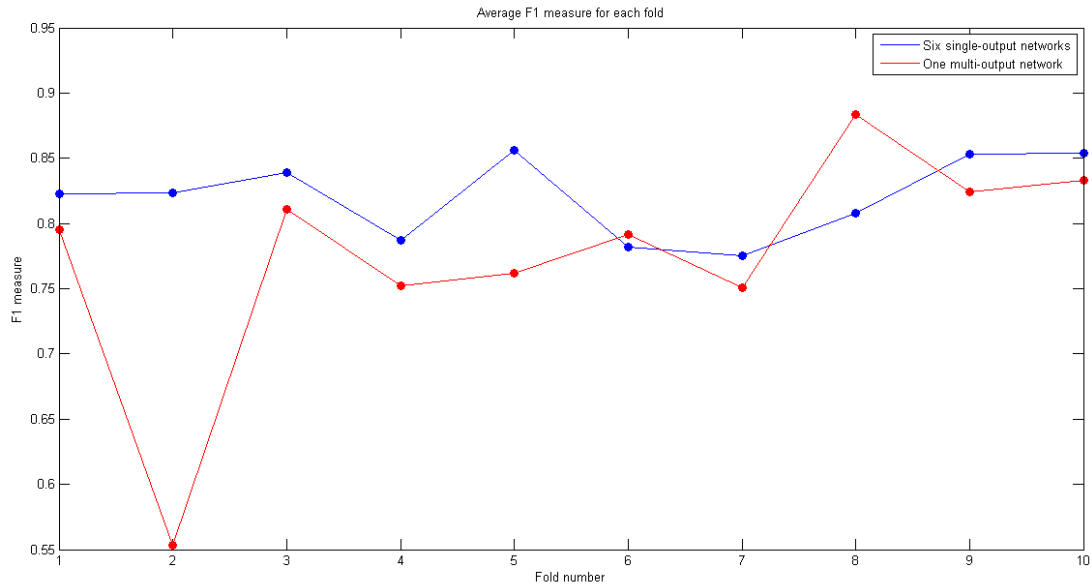


Figure 1: Average F_1 measure per fold for both types of networks

Clean dataset

Single six-output network

		Predicted class					
		1	2	3	4	5	6
Actual class	1	86	18	6	6	14	2
	2	9	167	6	5	11	0
	3	2	5	97	3	5	7
	4	0	6	1	203	4	2
	5	7	26	2	3	92	2
	6	0	3	12	2	2	188

Table 1: Confusion Matrix for single six-output ANN for the *clean* dataset

		Recall	Precision	F_1
Actual class	1	65%	83%	73%
	2	84%	74%	79%
	3	82%	78%	80%
	4	94%	91%	93%
	5	70%	72%	71%
	6	91%	94%	92%

Table 2: Recall, precision and F_1 measure for single six-output ANN for the *clean* dataset

$$C = \frac{833}{1004} = 83.0\%$$

Figure 2: Classification rate for single six-output ANN for the *clean* dataset

Six single-output networks

		Predicted class					
		1	2	3	4	5	6
Actual class	1	99	12	7	3	10	1
	2	9	168	3	6	11	1
	3	4	2	98	2	3	10
	4	2	4	2	205	1	2
	5	8	21	7	4	89	3
	6	0	3	10	5	2	187

Table 3: Confusion Matrix for six single-output ANNs for the *clean* dataset

		Recall	Precision	F_1
Actual class	1	75%	81%	78%
	2	84%	80%	82%
	3	82%	77%	80%
	4	95%	91%	93%
	5	67%	77%	72%
	6	90%	92%	91%

Table 4: Recall, precision and F_1 measure for six single-output ANNs for the *clean* dataset

$$C = \frac{846}{1004} = 84.3\%$$

Figure 3: Classification rate for six single-output ANNs for the *clean* dataset

Noisy dataset

Single six-output network

		Predicted class					
		1	2	3	4	5	6
Actual class	1	11	16	22	7	25	7
	2	4	149	18	6	6	4
	3	2	15	132	10	9	19
	4	3	7	13	175	4	7
	5	7	13	14	5	56	15
	6	1	4	15	5	9	186

Table 5: Confusion Matrix for single six-output ANN for the *noisy* dataset

		Recall	Precision	F_1
Actual class	1	13%	39%	19%
	2	80%	73%	76%
	3	71%	62%	66%
	4	84%	84%	84%
	5	51%	51%	51%
	6	85%	78%	81%

Table 6: Recall, precision and F_1 measure for single six-output ANN for the *noisy* dataset

$$C = \frac{709}{1001} = 70.8\%$$

Figure 4: Classification rate for single six-output ANN for the *noisy* dataset

Six single-output networks

		Predicted class					
		1	2	3	4	5	6
Actual class	1	15	11	25	8	23	6
	2	3	154	15	6	7	2
	3	1	11	139	8	11	17
	4	1	9	11	176	6	6
	5	4	8	19	4	67	8
	6	1	3	13	4	9	190

Table 7: Confusion Matrix for six single-output ANNs for the *noisy* dataset

		Recall	Precision	F_1
Actual class	1	17%	60%	27%
	2	82%	79%	80%
	3	74%	63%	68%
	4	84%	85%	85%
	5	61%	54%	58%
	6	86%	83%	85%

Table 8: Recall, precision and F_1 measure for six single-output ANNs for the *noisy* dataset

$$C = \frac{741}{1001} = 74.0\%$$

Figure 5: Classification rate for six single-output ANNs for the *noisy* dataset

Discussion of results

Overall, we can say that our neural networks perform quite well, especially on the *clean* dataset. There is one emotion in the *noisy* dataset however, for which the performance is unsatisfactory - emotion 1 with 17% recall and 60% precision for the six single-output networks. It is also worth noting, that the neural networks perform better than the decision trees on both *clean* and *noisy* datasets. The best classification rate for the *clean* dataset achieved by the decision trees was 74.0% compared to 84.3% scored by the neural networks. Similarly, for the *noisy* dataset the results were 61.9% compared to 74.0%.

Questions

Optimal topology

The final, optimal parameters for the two types of networks are:

1. Single six-output network:
 - Hidden layers: 1
 - Hidden neurons: 23
 - Transfer function: 'tansig' (sigmoid)
 - Training function: 'trainscg'
 - Learning rate: 0.01
 - Minimum gradient: $3 * 10^{-6}$
 - Goal: $2 * 10^{-3}$
2. Six single-output networks:
 - Hidden layers: 1
 - Hidden neurons: 9
 - Transfer function: 'tansig' (sigmoid)
 - Training function: 'trainscg'
 - Learning rate: 0.01
 - Minimum gradient: $5 * 10^{-6}$
 - Goal: $3 * 10^{-3}$

We arrived at these values by performing a series of simulations of the performance of networks with different values on the first fold on the *clean* dataset. In the first experiment, we tried to optimise the topology i.e. number of hidden layers, number of hidden neurons in each layer as well as the training function. We did this by repeatedly training the network with different parameters and plotted the results which are presented in **Figures 6, 7, 8, 9**. From the analysis of the results it can be concluded that the networks with just one hidden layer perform slightly better than those with two hidden layers, and so we trained our final networks with a single hidden layer. In case of the single six-output network, the peak is achieved for 32 hidden neurons, however as our final topology we selected 23 hidden neurons, since the classification rate for this topology is insignificantly smaller, yet the number of neurons is reduced by almost a third. For the six single-output networks, we decided to choose 9 hidden neurons.

In terms of the training function, it looks that 'trainscg' and 'trainrp' perform similarly, while 'trainlm' almost always exhibits the worst performance. Due to the speed of training and the classification rates at our chosen numbers of hidden neurons, we ultimately choose 'trainscg' as our training function for both types of networks. We also tried to look at the performance of 'trainbr' function, however its running time was extremely slow and we ended up not including it.

We selected these values because we needed eventually to decide on the topology of the network. However, by looking at the plots, it is clear that the number of hidden layers and neurons (at least for the range of values which we investigated) does not cause the performance to vary significantly beyond some minimum number of neurons and before some maximum. As long as the number of hidden neurons is in some reasonable range (above 20 for the single six-output network and between 6 and 30 for six single-output networks), it does not matter tremendously what exact value we would choose.

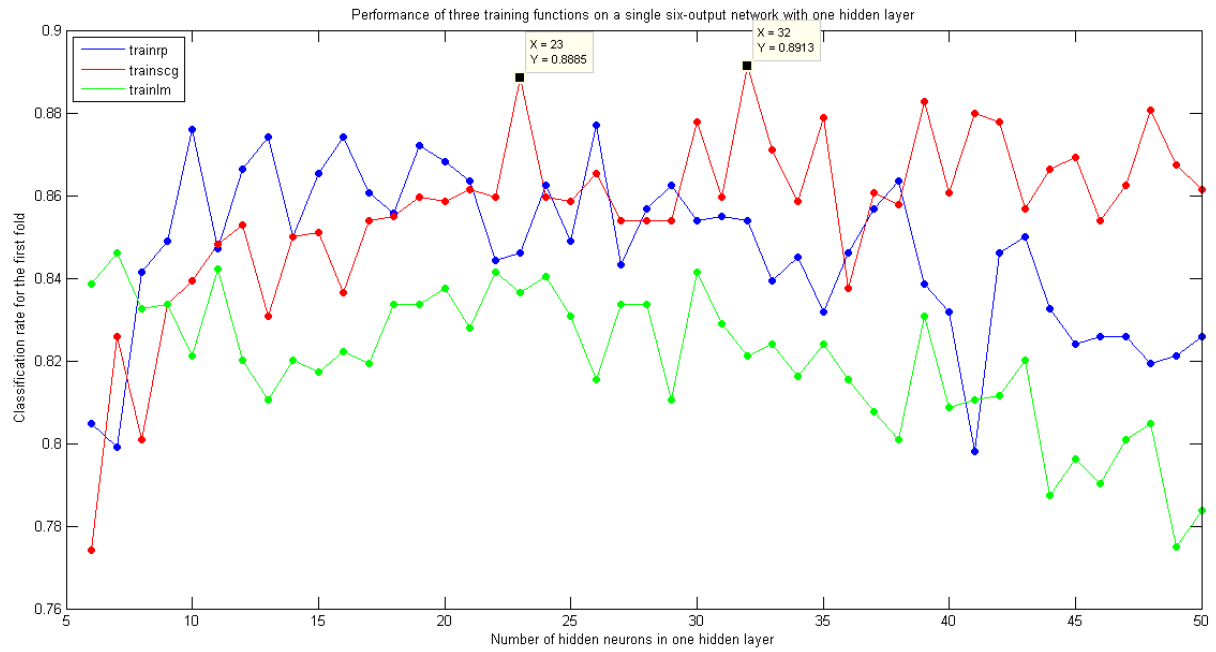


Figure 6: Classification rate for the first fold of the *clean* dataset with respect to the training function and the number of hidden neurons for a single six-output network with one hidden layer

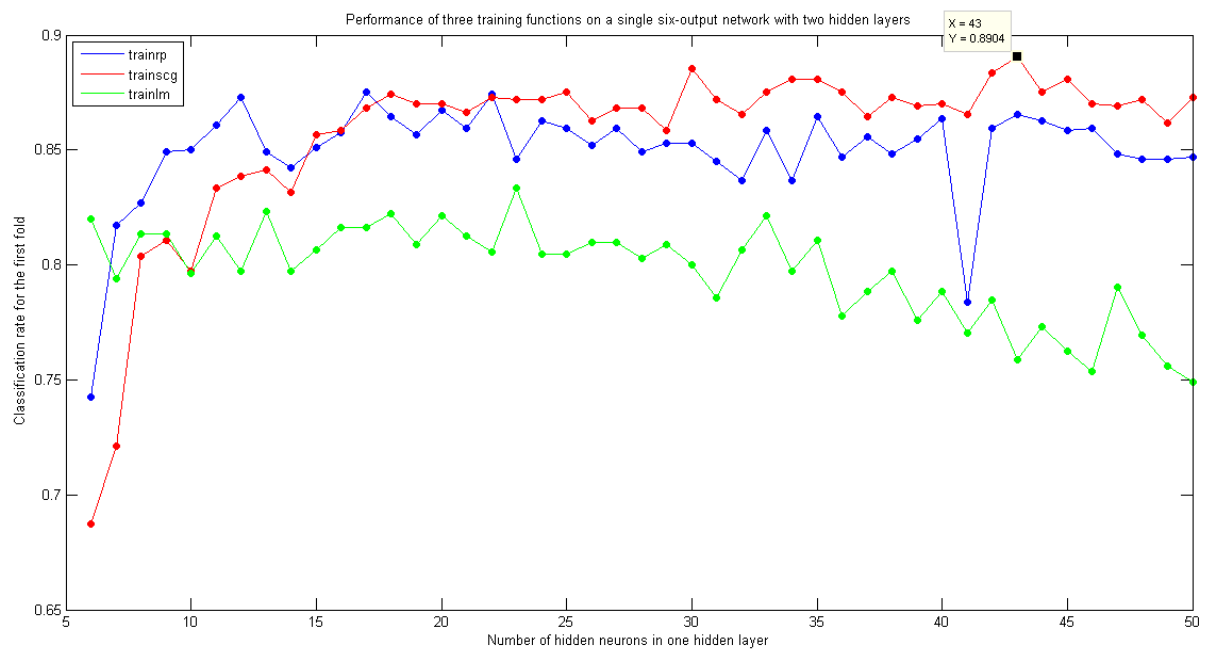


Figure 7: Classification rate for the first fold of the *clean* dataset with respect to the training function and the number of hidden neurons for a single six-output network with two hidden layers

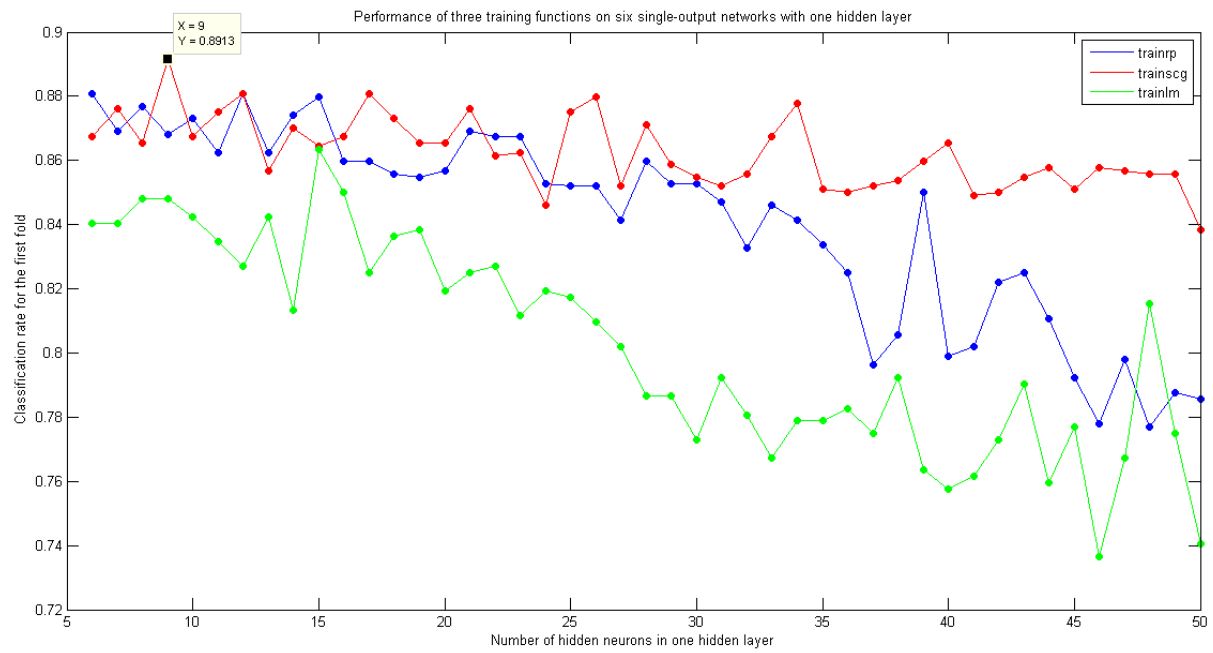


Figure 8: Classification rate for the first fold of the *clean* dataset with respect to the training function and the number of hidden neurons for six single-output networks with one hidden layer

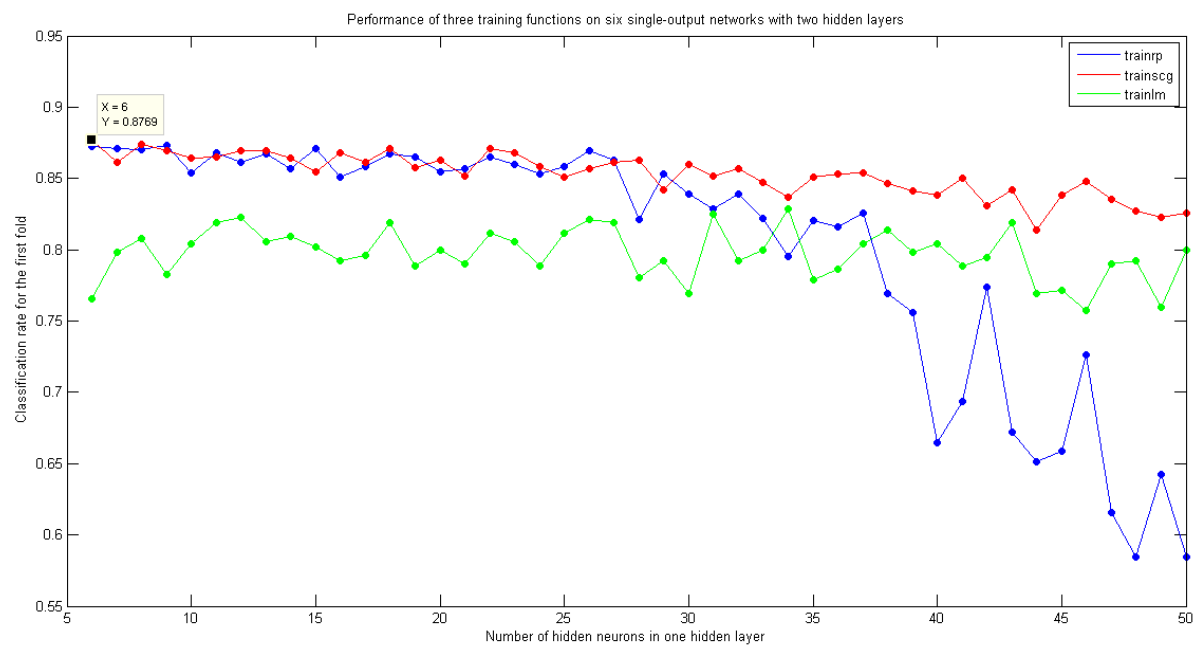


Figure 9: Classification rate for the first fold of the *clean* dataset with respect to the training function and the number of hidden neurons for six single-output networks with two hidden layers

Overfitting

There are several ways of avoid overfitting and therefore improve generalisation. General recommendation is to train network which is just large enough to provide a good fit. As mentioned in the book 'Pattern Recognition', excessive training on multi-layer network may result in poor generalisation 'as the net implements a complex decision boundary "tuned" to the specific training data'. In our experiment we focussed on networks which were as small as possible.

Another way to avoid overfitting is regularisation, which involves modifying performance function. Default performance function for 'trainscg' was mean squared error. We noticed that by changing our function to the mean squared error and the mean squared weight and bias 'msereg' we were able get results, which gave us higher classification rate, up to 1%. In order to further improve our results we decided to work on early stopping. This was done by experimenting with the goal and minimum gradient. We ran simulations for both types of networks during which we varied the two parameters and the results of these simulations are presented in Figures 10 and 11. It follows directly from the Figures that the optimal goal and minimum gradient for the six-output network are $2 * 10^{-3}$ and $3 * 10^{-6}$ respectively. For the six single-output networks it is $3 * 10^{-3}$ and $5 * 10^{-6}$.

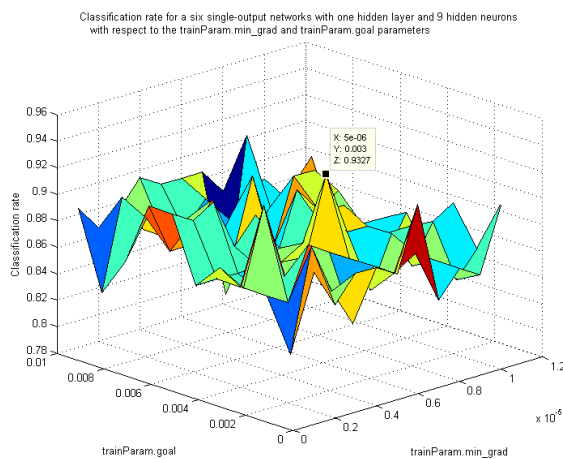


Figure 10: Classification rate for the first fold of the *clean* dataset with respect to the minimum gradient and the performance goal for six single-output networks with the optimal topology outlined before




Figure 11: Classification rate for the first fold of the *clean* dataset with respect to the minimum gradient and the performance goal for a single six-output network with the optimal topology outlined before

Six-output vs. single-output networks

The aggregated classification rate for the both types of networks is presented in Table 9.

	Classification rate <i>clean</i> dataset	Classification rate <i>noisy</i> dataset
Single six-output network	83.0%	70.8%
Six single-output networks	84.3%	74.0%

Table 9: Aggregated classification rates for all networks. Highlighted is the best result for a given dataset.

Clearly, the six single-output networks perform better in case of both datasets, however their advantage is not tremendous. This can be explained by the fact that six separate networks can be trained more specifically to detect one particular emotion than one six-output network. Just like in the real world, a General Practitioner can probably diagnose most common illnesses, yet six different specialist doctors might

achieve a smaller misdiagnosis rate. Additionally, six single-output networks can be more robust to noise, since during training, noise appearing in the emotions which the network is not concerned with have smaller effect (the class associated with them is always 0).

The main advantage of the single-output networks is therefore improved classification rate. Their disadvantages include a slower performance and slower training as well as the possibility of two networks classifying the same example as two different emotions. In our case, since we are using the sigmoid as the transfer function, in order to break tie we always selected the network with greatest output. Similarly, the single six-output network's main advantage is its speed which comes at the expense of the average classification rate.

Ideal parameter optimisation

We believe that for the n-fold cross-validation the parameters should be optimised for each fold separately. For every set of parameters and fold we should also record the performance measure. Then, the ideal way would be to choose a set of parameters which result in the largest average classification rate across all folds.

Another good way to optimise the parameters for the network trained on the entire *clean* dataset would be to discard the n-fold validation, and instead optimise the parameters on a single partition into training examples (probably around 900) and validation (the remaining examples). We do not need the test examples in this case, since all unseen data serves as such (including the one on which the network will be tested while marking this assignment).

Code Flowchart

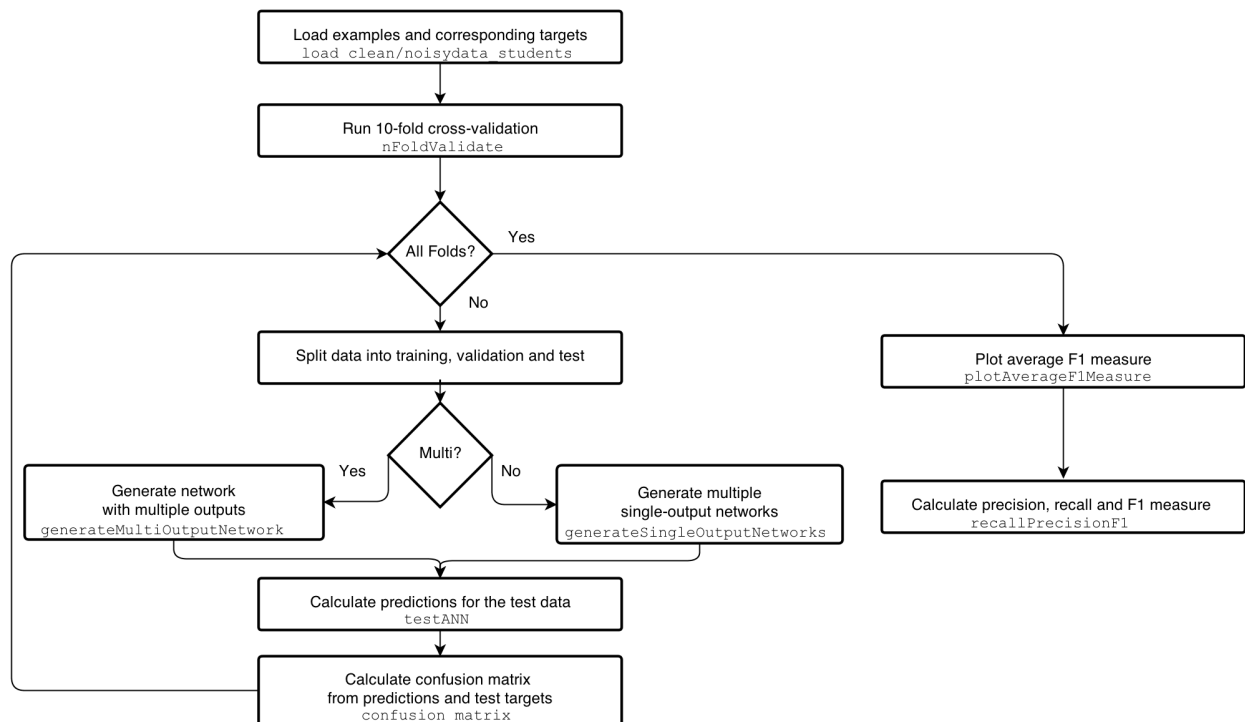


Figure 12: Code flowchart

To perform the 10-fold cross-validation and plot the average F_1 measure for each fold (note that due to some random variation the result will be different from the one presented in this report):

```
>> load cleandata_students.mat
>> singleOutputConfusionMatrix = nFoldValidate(x,y,10,'single');
>> multiOutputConfusionMatrix = nFoldValidate(x,y,10,'multi');
>> plotAverageF1Measure(singleOutputConfusionMatrix,multiOutputConfusionMatrix);
```

To generate a vector of predictions using the provided network:

```
>> predictions = testANN(net,x); % 'net' can be either type of network, 'x'
    represents the examples
```